

14_Optimizing_Loops_with_Numba

January 20, 2016

Text provided under a Creative Commons Attribution license, CC-BY. All code is made available under the FSF-approved MIT license. (c) Lorena A. Barba, 2013. Thanks: Gilbert Forsyth for help writing the notebooks. NSF for support via CAREER award 1149784. [LorenaABarba](https://twitter.com/LorenaABarba)

This notebook complements the [interactive CFD online](#) module **12 steps to Navier-Stokes**, addressing the issue of high performance with Python.

0.1 Optimizing Loops with Numba

You will recall from our exploration of [array operations with NumPy](#) that there are large speed gains to be had from implementing our discretizations using NumPy-optimized array operations instead of many nested loops.

[Numba](#) is a tool that offers another approach to optimizing our Python code. Numba is a library for Python which turns Python functions into C-style compiled functions using LLVM. Depending on the original code and the size of the problem, Numba can provide a significant speedup over NumPy optimized code.

Let's revisit the 2D Laplace Equation:

```
In [27]: from mpl_toolkits.mplot3d import Axes3D
         from matplotlib import cm
         import matplotlib.pyplot as plt
         import numpy as np

         ##variable declarations
         nx = 81
         ny = 81
         c = 1
         dx = 2.0/(nx-1)
         dy = 2.0/(ny-1)

         ##initial conditions
         p = np.zeros((ny,nx)) ##create a XxY vector of 0's

         ##plotting aids
         x = np.linspace(0,2,nx)
         y = np.linspace(0,1,ny)

         ##boundary conditions
         p[:,0] = 0                ##p = 0 @ x = 0
         p[:,-1] = y               ##p = y @ x = 2
         p[0,:] = p[1,:]           ##dp/dy = 0 @ y = 0
         p[-1,:] = p[-2,:]         ##dp/dy = 0 @ y = 1
```

Here is the function for iterating over the Laplace Equation that we wrote in Step 9:

```
In [17]: def laplace2d(p, y, dx, dy, l1norm_target):
    l1norm = 1
    pn = np.empty_like(p)

    while l1norm > l1norm_target:
        pn = p.copy()
        p[1:-1,1:-1] = (dy**2*(pn[2:,1:-1]+pn[0:-2,1:-1])+dx**2*(pn[1:-1,2:]+pn[1:-1,0:-2]))/(2*(dx**2+dy**2))
        p[0,0] = (dy**2*(pn[1,0]+pn[-1,0])+dx**2*(pn[0,1]+pn[0,-1]))/(2*(dx**2+dy**2))
        p[-1,-1] = (dy**2*(pn[0,-1]+pn[-2,-1])+dx**2*(pn[-1,0]+pn[-1,-2]))/(2*(dx**2+dy**2))

        p[:,0] = 0                ##p = 0 @ x = 0
        p[:,-1] = y                ##p = y @ x = 2
        p[0,:] = p[1,:]            ##dp/dy = 0 @ y = 0
        p[-1,:] = p[-2,:]          ##dp/dy = 0 @ y = 1
        l1norm = (np.sum(np.abs(p[:])-np.abs(pn[:])))/np.sum(np.abs(pn[:]))

    return p
```

Let's use the %%timeit cell-magic to see how fast it runs:

```
In [28]: %%timeit
    laplace2d(p, y, dx, dy, .00001)
```

1 loops, best of 3: 206 us per loop

Ok! Our function laplace2d takes around 206 micro-seconds to complete. That's pretty fast and we have our array operations to thank for that. Let's take a look at how long it takes using a more 'vanilla' Python version.

```
In [29]: def laplace2d_vanilla(p, y, dx, dy, l1norm_target):
    l1norm = 1
    pn = np.empty_like(p)
    nx, ny = len(y), len(y)

    while l1norm > l1norm_target:
        pn = p.copy()

        for i in range(1, nx-1):
            for j in range(1, ny-1):
                p[i,j] = (dy**2*(pn[i+1,j]+pn[i-1,j])+dx**2*(pn[i,j+1]-pn[i,j-1]))/(2*(dx**2+dy**2))

        p[0,0] = (dy**2*(pn[1,0]+pn[-1,0])+dx**2*(pn[0,1]+pn[0,-1]))/(2*(dx**2+dy**2))
        p[-1,-1] = (dy**2*(pn[0,-1]+pn[-2,-1])+dx**2*(pn[-1,0]+pn[-1,-2]))/(2*(dx**2+dy**2))

        p[:,0] = 0                ##p = 0 @ x = 0
        p[:,-1] = y                ##p = y @ x = 2
        p[0,:] = p[1,:]            ##dp/dy = 0 @ y = 0
        p[-1,:] = p[-2,:]          ##dp/dy = 0 @ y = 1
        l1norm = (np.sum(np.abs(p[:])-np.abs(pn[:])))/np.sum(np.abs(pn[:]))

    return p
```

```
In [30]: %%timeit
    laplace2d_vanilla(p, y, dx, dy, .00001)
```

10 loops, best of 3: 32 ms per loop

The simple Python version takes 32 milli-seconds to complete. Let's calculate the speedup we gained in using array operations:

```
In [35]: 32*1e-3/(206*1e-6)
```

```
Out[35]: 155.33980582524273
```

So NumPy gives us a 155x speed increase over regular Python code. That said, sometimes implementing our discretizations in array operations can be a little bit tricky.

Let's see what Numba can do. We'll start by importing the special function decorator `autojit` from the `numba` library:

```
In [36]: from numba import autojit
```

To integrate Numba with our existing function, all we have to do is prepend the `@autojit` function decorator before our `def` statement:

```
In [38]: @autojit
def laplace2d_numba(p, y, dx, dy, llnorm_target):
    llnorm = 1
    pn = np.empty_like(p)

    while llnorm > llnorm_target:
        pn = p.copy()
        p[1:-1,1:-1] = (dy**2*(pn[2:,1:-1]+pn[0:-2,1:-1])+dx**2*(pn[1:-1,2:]+pn[1:-1,0:-2]))/(
        p[0,0] = (dy**2*(pn[1,0]+pn[-1,0])+dx**2*(pn[0,1]+pn[0,-1]))/(2*(dx**2+dy**2))
        p[-1,-1] = (dy**2*(pn[0,-1]+pn[-2,-1])+dx**2*(pn[-1,0]+pn[-1,-2]))/(2*(dx**2+dy**2))

        p[:,0] = 0                                ##p = 0 @ x = 0
        p[:,-1] = y                                ##p = y @ x = 2
        p[0,:] = p[1,:]                            ##dp/dy = 0 @ y = 0
        p[-1,:] = p[-2,:]                          ##dp/dy = 0 @ y = 1
        llnorm = (np.sum(np.abs(p[:])-np.abs(pn[:])))/np.sum(np.abs(pn[:]))

    return p
```

The only lines that have changed are the `@autojit` line and also the function name, which has been changed so we can compare performance. Now let's see what happens:

```
In [39]: %%timeit
laplace2d_numba(p, y, dx, dy, .00001)
```

1 loops, best of 3: 137 us per loop

Ok! So it's not a 155x speed increase like we saw between vanilla Python and NumPy, but it is a non-trivial gain in performance time, especially given how easy it was to implement. Another cool feature of Numba is that you can use the `@autojit` decorator on non-array operation functions, too. Let's try adding it onto our vanilla version:

```
In [41]: @autojit
def laplace2d_vanilla_numba(p, y, dx, dy, llnorm_target):
    llnorm = 1
    pn = np.empty_like(p)
    nx, ny = len(y), len(y)
```

```

while l1norm > l1norm_target:
    pn = p.copy()

    for i in range(1, nx-1):
        for j in range(1, ny-1):
            p[i,j] = (dy**2*(pn[i+1,j]+pn[i-1,j])+dx**2*(pn[i,j+1]-pn[i,j-1]))/(2*(dx**2+dy**2))

p[0,0] = (dy**2*(pn[1,0]+pn[-1,0])+dx**2*(pn[0,1]+pn[0,-1]))/(2*(dx**2+dy**2))
p[-1,-1] = (dy**2*(pn[0,-1]+pn[-2,-1])+dx**2*(pn[-1,0]+pn[-1,-2]))/(2*(dx**2+dy**2))

p[:,0] = 0                ##p = 0 @ x = 0
p[:,-1] = y               ##p = y @ x = 2
p[0,:] = p[1,:]           ##dp/dy = 0 @ y = 0
p[-1,:] = p[-2,:]         ##dp/dy = 0 @ y = 1
l1norm = (np.sum(np.abs(p[:])-np.abs(pn[:])))/np.sum(np.abs(pn[:]))

return p

```

```

In [42]: %%timeit
         laplace2d_vanilla_numba(p, y, dx, dy, .00001)

```

1 loops, best of 3: 561 us per loop

561 micro-seconds. That's not quite the 155x increase we saw with NumPy, but it's close. And all we did was add one line of code.

So we have:

Vanilla Python: 32 milliseconds

NumPy Python: 206 microseconds

Vanilla + Numba: 561 microseconds

NumPy + Numba: 137 microseconds

Clearly the NumPy + Numba combination is the fastest, but the ability to quickly optimize code with nested loops can also come in very handy in certain applications.

```

In [ ]:

```

```

In [ ]:

```

```

In [42]:

```

```

In [1]: from IPython.core.display import HTML
        def css_styling():
            styles = open("../styles/custom.css", "r").read()
            return HTML(styles)
        css_styling()

```

```

Out[1]: <IPython.core.display.HTML at 0x36fbb10>

```

(The cell above executes the style for this notebook. We modified a style we found on the GitHub of [CamDavidsonPilon](https://github.com/CamDavidsonPilon), [[@Cmrn_DP](https://twitter.com/cmnrn_dp)](https://twitter.com/cmnrn_dp).)